Weak-Signal Interpretation for Perimeter Anomaly Triage

Toward a Stratified, Human-Review-Gated Architecture for Early Perimeter Detection

Stephen A. Putman
putmanmodel@pm.me

Version 0.2 — April 2026

Abstract

Perimeter detection systems often begin not with cleanly resolved events, but with weak, noisy, partial, and locally ambiguous irregularities: slight motion persistence near a boundary, weak thermal inconsistency, intermittent fence-line disturbance, unusual route crossing, acoustic irregularity, low-confidence object-track deviation, or disagreement across sensing channels. In such cases, the architectural problem is not only detection, but promotion control: how weak findings should influence attention before they are allowed to become stronger operational significance. This bridge paper argues that perimeter anomaly triage is a natural application domain for a stratified weak-signal architecture organized around four bounded stages: bounded findings, provisional candidate meanings, short-horizon attentional escalation, and governed promotion under explicit human review. The paper contributes a minimal transfer surface for that claim: a light state ladder, a minimal event object, illustrative promotion predicates, and compact examples showing how weak anomalies may matter before they justify stronger consequence. It does not propose autonomous enforcement, identity resolution, or a full surveillance platform. Its narrower contribution is architectural: perimeter systems need a disciplined place for weak anomaly signals to matter before they harden into alerts, incident states, or ignored noise.

1. Introduction

Perimeter detection systems rarely begin with certainty. More often, they first encounter weak, noisy, partial, and locally ambiguous irregularities: a slight motion persistence near a fence line, a weak thermal inconsistency, an unusual route crossing near a boundary, intermittent acoustic disturbance, low-confidence object-track irregularity, or disagreement across sensing channels. Any one of these may be benign. Any one of them may also be the first visible sign that ordinary background continuity is no longer fully trustworthy.

That creates an architectural problem. If the system ignores weak but structured departures because they do not yet justify confidence, it loses useful early warning. If it commits too early, it converts provisional significance into brittle alerting, unnecessary review burden, or stronger consequence unsupported by the evidence. The relevant question is therefore not only whether an anomaly exists, but how a perimeter system should behave when evidence is incomplete, weakly convergent, and still locally ambiguous. In operational terms, the middle state is intended to reduce premature hard alerts, preserve early warning lead time, and keep review load bounded by requiring stronger consequence to pass through explicit promotion gates.

This paper takes a narrow position. Perimeter anomaly triage is a natural application domain for a stratified weak-signal architecture that separates bounded observation, provisional interpretation, short-horizon attentional escalation, and governed consequence under explicit human-review gating. The point is disciplined escalation under uncertainty in systems that must remain sensitive to weak irregularities without allowing them to harden prematurely into incident state.

The contribution is architectural rather than algorithmic. It argues for a minimal middle layer in which weak anomaly signals can raise attention before they are allowed to become stronger operational consequence.

2. Why Perimeter Detection Needs a Middle State

Perimeter systems operate at the boundary between ordinary background variation and possible incursion, interference, or sensor irregularity. That matters because many meaningful anomalies first appear not as resolved events, but as partial disturbances distributed across several bounded channels. A motion system may show weak persistence while thermal evidence remains inconclusive. A fence-line disturbance may recur without yielding a clean object track. A route anomaly may appear while camera confidence drops or acoustic evidence becomes ambiguous. In each case, the system may have enough structure to justify heightened attention, but not enough to justify durable conclusion.

Without an explicit middle state, architectures tend to collapse in one of two directions. In the first, weak anomalies are treated as ordinary variance until they become obvious, by which point the value of early handling has been reduced. In the second, weak irregularities are promoted too quickly into hard alerts, stronger review burden, or forced incident framing, producing brittleness and false certainty. A better architecture preserves provisional significance without premature commitment.

The key point is simple: perimeter systems need somewhere for weak anomaly signals to live before they harden into alerts, incidents, or ignored noise.

That middle state is useful because early perimeter handling is not only a recognition problem. It is also a promotion-control problem. The system must decide what weak findings are allowed to influence next-step attention, what candidate meanings may remain live under ambiguity, what forms of temporary watchfulness are permitted, and what stronger outcomes require persistence, convergence, corroboration, and explicit review. Without those distinctions, a perimeter stack becomes either dull or jumpy. Neither is trustworthy.

This paper therefore treats perimeter anomaly handling as an architectural transfer problem. The same broader weak-signal logic already established elsewhere applies here: weak findings may matter before they settle meaning, repeated or convergent evidence may widen scope, and stronger consequence should require more than short-horizon attention alone. In perimeter terms,

the system needs a bounded place for provisional significance before stronger operational consequence becomes eligible.

3. Core Claim and Design Principle

The central claim of this paper is simple: bounded findings should raise attention first, candidate meanings should remain provisional, and stronger perimeter consequence should occur only through persistence, convergence, rationale, and explicit human-review-gated promotion. Promotion control, in this paper, refers to the declared rules and review boundary governing when weak findings may move from provisional significance into stronger triage consequence.

This framing does not reject sensing, classification, or ordinary perimeter monitoring practice. It reframes their early-stage role. In many perimeter settings, the first useful evidence is too partial to deserve strong commitment but too structured to deserve indifference. The architecture should therefore preserve a strict separation between bounded observation, provisional interpretation, short-horizon attentional escalation, and governed consequence.

The governing design principle is equally simple: weak perimeter anomalies should influence attention before they justify stronger operational consequence.

That means a slight motion persistence, a weak thermal inconsistency, an intermittent fence-line disturbance, or a recurring route irregularity can widen attention, tighten inspection, or trigger bounded corroboration behavior without immediately becoming durable incident state, strong target belief, or higher-impact response. This is not a claim that weak signals are secretly decisive. It is a claim that they are often useful earlier than hard operational thresholds allow, and that a disciplined architecture should preserve that usefulness while still demanding stronger support for stronger consequence.

This paper is an argument for promotion control under uncertainty. Weak findings may matter provisionally. Repeated or convergent evidence may widen scope. But stronger consequence should require more than short-horizon attention alone.

4. Stratified Architecture Mapping

The four layers described below are processing functions. The ladder introduced later describes the lifecycle state of a case as those functions operate on it.

The proposed architecture separates perimeter anomaly handling into four stages. This is not an argument against cameras, thermal sensing, motion analytics, route analysis, acoustic monitoring, or cross-sensor fusion. It is an argument for preserving architectural boundaries around what each stage is allowed to do. In perimeter settings, early irregularities often arrive in weak, partial, or cross-channel form. If every stage is allowed to behave as though its local signal were already a resolved incident, the system becomes brittle. A more disciplined architecture

separates bounded observation, provisional interpretation, short-horizon attention shaping, and governed promotion.

4.1 Bounded Observation

At the first layer, the system emits narrow findings only. These findings remain local to their sensing pathway or analysis channel. They may include weak motion persistence, slight thermal inconsistency, intermittent fence-line disturbance, low-confidence object-track irregularity, unusual route crossing near a boundary, acoustic anomaly, lighting or camera inconsistency, or disagreement across sensing channels.

At this stage, the system is not yet claiming intruder, threat, trespass, or incident. It is preserving local structure in a bounded form so that downstream layers can compare, combine, and challenge it. Observation remains narrow specifically so that one local subsystem does not get to define the whole situation prematurely.

4.2 Provisional Interpretation

At the second layer, the system forms candidate local meanings from one or more bounded findings. These are not final truths. They are admissible interpretations of a partial pattern. Examples include possible human presence, possible animal movement, possible environmental disturbance, possible sensor fault, possible route anomaly, or unresolved perimeter irregularity needing more evidence.

The crucial requirement is that ambiguity remain visible. The system does not become more trustworthy by pretending uncertainty has vanished. It becomes more trustworthy by representing candidate meanings in a way that remains revisable. Competing interpretations may coexist. A perimeter stack should be able to say, in effect, that the current pattern is worth attention without pretending it has already been settled.

4.3 Short-Horizon Attention

At the third layer, weak but informative deviation influences short-horizon attention. Reflex salience may be used as a local mechanism within this layer, but the architectural role remains short-horizon attention shaping rather than persistence or stronger consequence. This is the stage at which the system may tighten an observation window, increase local sampling, cue another sensor, preserve a temporary watch state, raise local review priority, or slow commitment while additional evidence is gathered.

The important constraint is that this layer remains temporary and non-authorizing. It shapes what the system inspects next, but it does not by itself write durable memory, mutate long-horizon baselines, resolve identity, or authorize stronger consequence. Short-horizon attention is useful precisely because it gives the system a way to respond to weak significance without forcing structural commitment.

4.4 Governed Promotion / Human Review

Only after persistence, convergence, or stronger support builds should the architecture permit governed promotion. At this stage, the system may allow escalation into a human-review queue, monitored watch-zone elevation, temporary quarantine of a suspect sensor stream, bounded reviewer alerting, or stronger incident-candidate routing. P4 is not a salience state. It is a reviewed promotion state, entered only when declared predicates are satisfied and the supporting evidence package is preserved for inspection.

The point is not that every anomaly must escalate. The point is that escalation should be explicit, inspectable, and predicate-bound rather than silently triggered by local salience alone. Human review matters here not as decorative oversight, but as a hard boundary against premature hardening. Stronger significance should require declared support, preserved rationale, and explicit review rather than automatic collapse from weak cue to stronger consequence.

A compact pipeline expresses the structure clearly:

local finding → candidate meaning → bounded attention shift → governed promotion

Or, in the broader stack language:

notice → interpret → shape → govern

5. Minimal Operational Contract

This section keeps the transfer concrete without turning the paper into a full surveillance platform specification. The point is not to prescribe one implementation. It is to define a minimal contract surface that makes the architecture inspectable and portable. A good bridge paper should show what kinds of states, records, and transitions are minimally required for the transfer to be meaningful.

5.1 Light state ladder

P-states are progression states. R0 and Q0 are terminal resolution outcomes.

A minimal state ladder is enough to distinguish weak-signal attention from stronger operational consequence:

- **P0 — Normal:** within the expected corridor for the current zone, boundary condition, and operating mode.
- **P1 — Noticed:** a weak local deviation is recorded as a bounded finding.
- **P2 — Watching:** the deviation persists across a short horizon, recurs within a bounded interval, or remains significant enough to justify additional sensing.

- **P3 — Corroborated:** a second supporting channel, repeated observation, or consistent local pattern supports the concern.
- **P4 — Escalated:** a governed review boundary has been crossed and the case is eligible for stronger triage significance under human review.
- **R0 — Resolved:** the deviation returns to corridor without stronger support.
- **Q0 — Quarantined:** evidence favors sensor artifact, environmental noise, or another non-security explanation requiring bounded isolation rather than stronger consequence.

The key boundary is simple: P1–P3 are attention and evidence states; P4 belongs to governed consequence.

5.2 Minimal event object

A weak-signal architecture needs a small event object that preserves structure without forcing immediate incident semantics. In minimal form, such an event should carry:

- event identity and time bounds
- zone or boundary scope
- source channel or channels
- observed signal values or local deltas
- current state in the anomaly ladder
- candidate meanings
- short-horizon actions taken
- supporting channels or corroborators
- confidence or uncertainty regime
- closure reason if resolved or quarantined

The point is not to freeze a universal schema. The point is to preserve enough context that weak deviations remain reviewable, comparable, and governable.

An illustrative event record might look like this in compact form:

{ id, first_seen, last_seen, zone, source_channels[], finding_values, state, candidate_meanings[], corroborators[], short_horizon_actions[], confidence_regime, closure_reason }

That level of structure is enough to make transitions visible. A reviewer or implementer can ask concrete questions: What entered the layer? What candidate meanings were live? What changed the state? What support was present? Why did the case resolve, quarantine, or escalate?

5.3 Transition logic and promotion predicates

The transition logic should remain explicit, even in a lightweight deployment.

- **P0 → P1:** a weak threshold is exceeded for the current zone or operating mode, and data-quality checks pass.
- **P1 → P2:** persistence, recurrence, or bounded continued significance is observed within a short review horizon.
- **P2 → P3:** convergence or corroboration appears across distinct channels, repeated observations, or compatible local patterns.
- **P3 → P4:** a governed escalation gate is met under declared review rules and preserved rationale.
- **P1/P2/P3 → R0:** the deviation returns to corridor and no corroboration arrives within the allowed horizon.
- **Any state → Q0:** evidence increasingly favors sensor fault, environmental disturbance, non-security explanation, or another bounded artifact requiring isolation rather than stronger consequence.

Promotion should require more than short-horizon attention alone. At minimum, stronger perimeter significance should be gated by declared predicates such as:

- persistence across a declared short review window
- corroboration from at least two independent findings or sensing paths within scope
- repeated reappearance of the same candidate pattern within scope
- a rationale record sufficient for later inspection
- explicit human review or another declared, auditable promotion gate for stronger triage significance

These are examples rather than fixed thresholds. Here, independence refers not merely to two outputs, but to distinct sensing paths, physical vantage points, or failure modes sufficient to reduce trivial self-corroboration.

The architectural requirement is that promotion be declared, inspectable, and predicate-bound rather than silently triggered by interest or urgency alone.

At minimum, a promotion record should preserve the candidate label, supporting findings, relevant zone or boundary scope, confidence regime, transition reason, and reviewer or gate identity. Promotion may be performed by an automated queueing gate, a human reviewer, or a hybrid review path, but in all cases the transition should preserve who or what authorized it, which predicates were satisfied, and what evidence supported the move. If those predicates are not met, the candidate remains provisional, resolves, or is quarantined.

6. Three Compact Examples

The following examples are deliberately compact. Their purpose is not to simulate a full perimeter platform, but to show how weak findings can influence attention before they justify stronger operational consequence. In each case, the point is the same: weak signals are not ignored, but they are also not allowed to harden prematurely into stronger incident significance.

6.1 Thermal + motion + route irregularity

Consider a perimeter segment in which three weak irregularities appear within a bounded horizon. First, a thermal channel records a faint moving inconsistency near the boundary. Second, motion persistence appears intermittently in roughly the same local region. Third, a route-tracking or path-analysis layer detects an unusual crossing pattern near the same zone relative to recent local conditions.

No single signal justifies a claim such as intrusion confirmed. The correct architectural response is therefore not immediate incident declaration. Instead, the system preserves the findings separately with scope, uncertainty, and provenance, then forms a provisional candidate state in which possible human presence and possible environmental disturbance may coexist as admissible meanings. Because the pattern is weakly convergent, short-horizon attention rises. The system may tighten the observation window, cue an additional sensor, increase local sampling, or preserve a temporary watch state for the affected zone.

If persistence builds or corroboration strengthens, the case may move toward governed escalation under explicit review predicates. If the pattern dissipates, fails to recur, or is better explained by a non-security condition, the case resolves or is quarantined without confusion. The key point is that faint convergence matters before certainty exists, but it matters first as attention, not as stronger consequence.

6.2 Fence-line disturbance with ambiguity

Consider a fence-line segment in which a small visual or vibration disturbance appears near the boundary, but camera confidence is inconsistent and no clean object track is formed. The disturbance is too structured to ignore, especially if it recurs, yet too incomplete to deserve a hard security event on first appearance.

At the bounded-observation layer, the system records only the local disturbance and its uncertainty. At the provisional-interpretation layer, several candidate meanings remain live: possible fence interaction, possible wind or vegetation disturbance, possible animal movement, or possible sensor artifact. The architecture does not become stronger by forcing these possibilities into one premature label.

Because the finding is locally meaningful, the system may enter a bounded attention state. It may widen the observation window, request additional sensing, preserve a temporary watch

condition, or increase local review priority for that fence segment. If the same pattern repeats or gains corroboration from another channel, human-review-gated escalation becomes more eligible. If confirming evidence fails to appear, the candidate decays, resolves, or is quarantined. In this example, the usefulness of the architecture lies in preserving the difference between “not nothing” and “not yet enough.”

6.3 Animal vs. human vs. sensor artifact ambiguity

Consider a localized anomaly in which motion, low-confidence thermal structure, and inconsistent visual evidence appear within a perimeter zone, but the available pattern could plausibly reflect animal movement, human presence, vegetation or weather disturbance, or sensor artifact. This kind of ambiguity is common in real perimeter conditions and is exactly where single-stage certainty performs badly.

A disciplined architecture preserves those candidate meanings explicitly rather than pretending one has already won. The system records the local findings, forms a provisional anomaly candidate, and uses short-horizon attention only to widen scope in bounded ways. Another sensor may be cued. The review horizon may be tightened. Local attention may rise. But stronger consequence does not follow from ambiguity alone.

Only if the pattern strengthens, persists, or converges with additional evidence should the system permit escalation into stronger triage significance under explicit human review. If it remains ambiguous but harmless, it may resolve without incident. If it increasingly resembles sensor artifact or environmental noise, it may be quarantined. The central lesson is simple: ambiguity is not failure. It is a state the architecture must be able to carry without collapsing into either indifference or false certainty.

7. Failure Modes and Boundary Discipline

A perimeter weak-signal architecture is useful only if it preserves its boundaries under pressure. Most failure modes arise when one layer is allowed to silently perform the job of another.

The first failure mode is premature hardening. A weak local cue is treated as though it already justifies incident-level significance. In practice, this means a single thermal irregularity, low-confidence motion trace, or fence-line disturbance is allowed to collapse directly into stronger operational framing. Such collapse reduces ambiguity visibly, but only by hiding it. The result is brittle escalation and avoidable false certainty.

The second failure mode is salience silently becoming persistence. Short-horizon attention exists to shape what the system inspects next, not to write durable state or stronger consequence on its own. If a temporary watch condition silently mutates into a stronger case record without declared predicates and preserved rationale, the architecture loses the very discipline it was meant to provide.

The third failure mode is forced single-label interpretation under ambiguity. A perimeter system may encounter patterns that genuinely support more than one admissible reading. If the architecture suppresses that competition too early, later correction becomes harder and trustworthiness declines. Ambiguity is not a weakness to be hidden; it is part of the object the system is handling.

The fourth failure mode is operator review being treated as decorative rather than structural. In this paper, human review is not a cosmetic final step. It is part of the promotion boundary. If stronger triage significance is allowed to emerge without explicit review or an equivalent governed gate, the architecture drifts toward automated hardening from weak evidence, which is precisely what it is meant to avoid.

The fifth failure mode is unbounded promotion of sensor artifact or environmental disturbance. Perimeter environments contain motion clutter, weather effects, vegetation movement, lighting changes, animal activity, and ordinary sensor irregularity. A trustworthy architecture must therefore preserve not only escalation paths, but also bounded resolution and quarantine paths. Systems that can escalate but cannot cleanly back down become noisy, distrusted, and operationally expensive.

These failure modes point toward a simple discipline. Observation should remain narrow. Interpretation should remain revisable. Short-horizon attention should remain temporary. Promotion should remain declared, reviewable, and predicate-bound. If those boundaries hold, weak anomaly handling becomes more legible, less brittle, and less prone to both indifference and overreaction.

8. Conclusion

Perimeter anomaly triage is a natural application domain for a stratified weak-signal architecture in which bounded findings remain local, candidate meanings remain provisional, short-horizon attention may rise without becoming persistence, and stronger consequence occurs only through persistence, convergence, rationale, and explicit human-review-gated promotion.

The contribution of this paper is narrow. It is a bridge note showing that perimeter systems need a disciplined place for weak anomaly signals to matter before they harden into alerts, incident states, or ignored noise.

That is the architectural transfer. Weak perimeter anomalies should influence attention before they justify stronger operational consequence. A system that can hold that middle ground is better equipped to remain sensitive to faint irregularities without allowing uncertainty to collapse prematurely into brittle certainty.